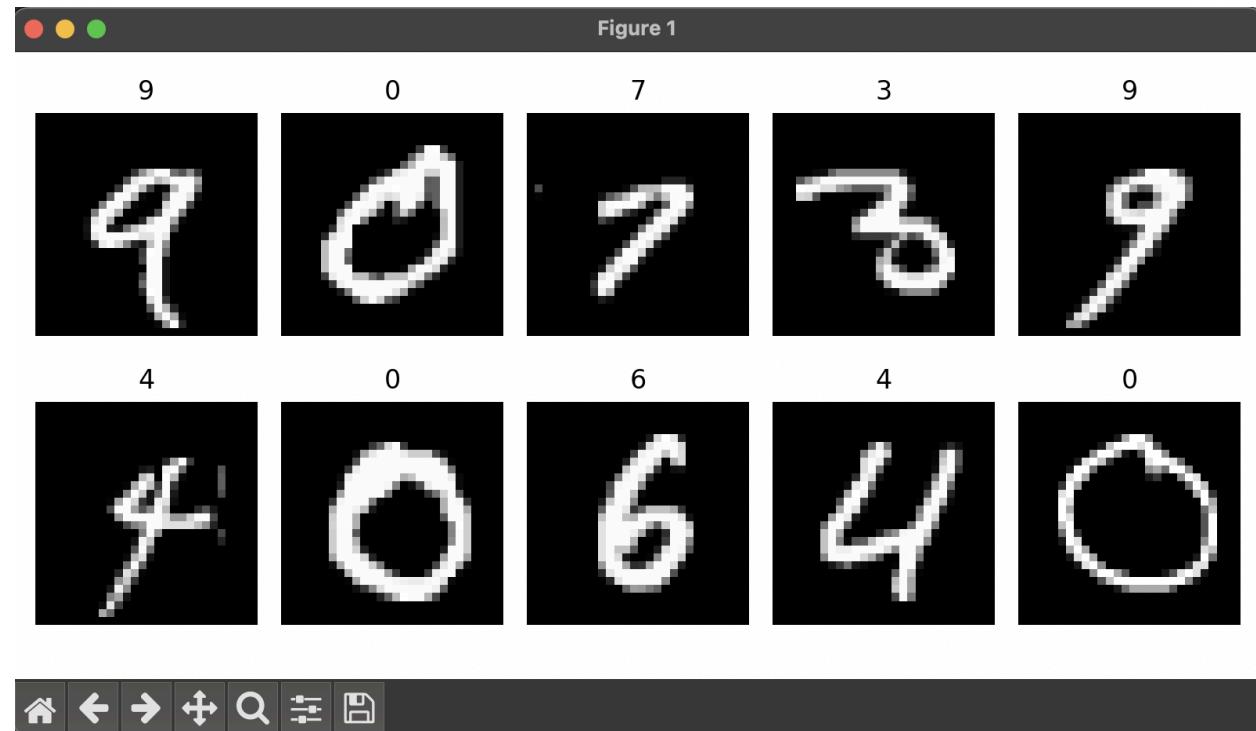


## Part A: MLP

### Task 1



### Task 2

```
plt.show()
Epoch 01 | Train Loss 2.4393 Acc 10.02% | Test Loss 2.4422 Acc 10.28%
Epoch 02 | Train Loss 2.4295 Acc 9.93% | Test Loss 2.4314 Acc 10.28%
Epoch 03 | Train Loss 2.4279 Acc 10.22% | Test Loss 2.4117 Acc 9.80%
Epoch 04 | Train Loss 2.4268 Acc 10.32% | Test Loss 2.3742 Acc 10.28%
Epoch 05 | Train Loss 2.4263 Acc 9.92% | Test Loss 2.3626 Acc 11.35%
(sharvaripurighalla@sharvari-mac EE456ComputerAssignment3 %)
```

The training and test accuracies remain around 10%, which is roughly the same as random guessing among ten digit classes (0-9). Loss values stay near 2.43 → 2.36, showing no meaningful learning trend.

This happens because:

- The MLP is extremely small ( $784 \rightarrow 8 \rightarrow 4 \rightarrow 10$ ) with very few neurons to capture complex digit patterns
- The learning rate (0.5) is too high, causing unstable or ineffective gradient updates.
- The only 5 epochs of training provide very limited optimization steps

Thus, this result is **not good**. The network is effectively guessing and has not learned to recognize digits. It's a baseline that confirms the model and training loop run correctly but requires larger capacity, lower learning rate, or more epochs to perform well.

### Task 3

```
from __future__ import print_function
import tensorflow as tf
import numpy as np
import random
import time
import os
import sys
import math
import matplotlib.pyplot as plt
plt.show()
Epoch 01 | Train Loss 2.4393 Acc 10.02% | Test Loss 2.4422 Acc 10.28%
Epoch 02 | Train Loss 2.4295 Acc 9.93% | Test Loss 2.4314 Acc 10.28%
Epoch 03 | Train Loss 2.4279 Acc 10.22% | Test Loss 2.4117 Acc 9.80%
Epoch 04 | Train Loss 2.4268 Acc 10.32% | Test Loss 2.3742 Acc 10.28%
Epoch 05 | Train Loss 2.4263 Acc 9.92% | Test Loss 2.3626 Acc 11.35%
Epoch 06 | Train Loss 2.4263 Acc 10.04% | Test Loss 2.4972 Acc 9.74%
Epoch 07 | Train Loss 2.4256 Acc 9.88% | Test Loss 2.5809 Acc 10.32%
Epoch 08 | Train Loss 2.4252 Acc 10.08% | Test Loss 2.4438 Acc 10.09%
Epoch 09 | Train Loss 2.4226 Acc 10.21% | Test Loss 2.3562 Acc 10.10%
Epoch 10 | Train Loss 2.4275 Acc 10.14% | Test Loss 2.3537 Acc 10.32%
Epoch 11 | Train Loss 2.4288 Acc 10.06% | Test Loss 2.3824 Acc 8.92%
Epoch 12 | Train Loss 2.4277 Acc 10.24% | Test Loss 2.4947 Acc 10.10%
Epoch 13 | Train Loss 2.4276 Acc 10.03% | Test Loss 2.3818 Acc 10.10%
Epoch 14 | Train Loss 2.4252 Acc 10.08% | Test Loss 2.4379 Acc 10.28%
Epoch 15 | Train Loss 2.4286 Acc 10.01% | Test Loss 2.4809 Acc 10.32%
> (ee456) sharvaripurighalla@sharvari-mac EE456ComputerAssignment3 %
```

Even after 15 epochs, both training and testing accuracy remain around 10%, which still reflects random guessing among the 10 digit classes. The loss value does not meaningfully decrease, suggesting the model is not learning any useful features. The network is still too shallow, providing insufficient representational capacity to capture MNIST digit patterns. The learning rate is also too high, and simply increasing the number of epochs does not help when the model capacity of learning rate are the true bottlenecks.

### Task 4

→ Batch size = 32:

```
E456ComputerAssignment3/mlp.py
Epoch 01 | Train Loss 2.4156 Acc 10.24% | Test Loss 2.3592 Acc 9.74%
Epoch 02 | Train Loss 2.3586 Acc 10.20% | Test Loss 2.4312 Acc 11.35%
Epoch 03 | Train Loss 2.3625 Acc 10.11% | Test Loss 2.3250 Acc 9.80%
Epoch 04 | Train Loss 2.3602 Acc 10.17% | Test Loss 2.3417 Acc 9.80%
Epoch 05 | Train Loss 2.3596 Acc 10.20% | Test Loss 2.3603 Acc 9.58%
Epoch 06 | Train Loss 2.3575 Acc 10.13% | Test Loss 2.3303 Acc 9.82%
Epoch 07 | Train Loss 2.3591 Acc 9.95% | Test Loss 2.3485 Acc 10.32%
Epoch 08 | Train Loss 2.3626 Acc 10.14% | Test Loss 2.3853 Acc 11.35%
Epoch 09 | Train Loss 2.3594 Acc 10.26% | Test Loss 2.3354 Acc 9.74%
Epoch 10 | Train Loss 2.3605 Acc 10.19% | Test Loss 2.3223 Acc 10.10%
Epoch 11 | Train Loss 2.3592 Acc 10.23% | Test Loss 2.3643 Acc 11.35%
Epoch 12 | Train Loss 2.3622 Acc 10.20% | Test Loss 2.3662 Acc 10.10%
Epoch 13 | Train Loss 2.3591 Acc 10.19% | Test Loss 2.3573 Acc 9.82%
Epoch 14 | Train Loss 2.3577 Acc 10.29% | Test Loss 2.3347 Acc 10.10%
Epoch 15 | Train Loss 2.3576 Acc 10.10% | Test Loss 2.3318 Acc 10.32%
```

→ Batch size = 64:

Epoch	01	Train Loss	2.3443	Acc	10.15%		Test Loss	2.3216	Acc	8.92%
Epoch	02	Train Loss	2.3409	Acc	10.26%		Test Loss	2.3589	Acc	11.35%
Epoch	03	Train Loss	2.3407	Acc	10.08%		Test Loss	2.3492	Acc	11.35%
Epoch	04	Train Loss	2.3422	Acc	10.25%		Test Loss	2.3328	Acc	9.80%
Epoch	05	Train Loss	2.3399	Acc	10.25%		Test Loss	2.3501	Acc	9.80%
Epoch	06	Train Loss	2.3395	Acc	10.20%		Test Loss	2.3457	Acc	10.28%
Epoch	07	Train Loss	2.3384	Acc	10.22%		Test Loss	2.3221	Acc	9.82%
Epoch	08	Train Loss	2.3426	Acc	10.03%		Test Loss	2.3290	Acc	9.74%
Epoch	09	Train Loss	2.3406	Acc	10.33%		Test Loss	2.3643	Acc	9.74%
Epoch	10	Train Loss	2.3405	Acc	10.22%		Test Loss	2.3168	Acc	11.35%
Epoch	11	Train Loss	2.3389	Acc	10.25%		Test Loss	2.3176	Acc	11.35%
Epoch	12	Train Loss	2.3444	Acc	10.27%		Test Loss	2.3412	Acc	9.74%
Epoch	13	Train Loss	2.3416	Acc	10.03%		Test Loss	2.3422	Acc	9.80%
Epoch	14	Train Loss	2.3394	Acc	10.22%		Test Loss	2.3536	Acc	10.10%
Epoch	15	Train Loss	2.3389	Acc	10.30%		Test Loss	2.3536	Acc	10.32%

Accuracy: both runs give ~10% accuracy, which is again near random chance, meaning the model still isn't learning meaningful patterns.

Speed: with batch = 64, each epoch completed slightly faster because the model performed fewer weight-update steps per epoch (fewer batches total)

Loss trend: loss values stayed nearly constant, confirming the learning rate is still too high or the model is too simple.

Thus, increasing the batch size did not improve accuracy but reduced training time per epoch, which matches my theoretical expectations.

## Task 5

→ 0.5 Learning Rate:

Epoch	01	Train Loss	2.3443	Acc	10.15%		Test Loss	2.3216	Acc	8.92%
Epoch	02	Train Loss	2.3409	Acc	10.26%		Test Loss	2.3589	Acc	11.35%
Epoch	03	Train Loss	2.3407	Acc	10.08%		Test Loss	2.3492	Acc	11.35%
Epoch	04	Train Loss	2.3422	Acc	10.25%		Test Loss	2.3328	Acc	9.80%
Epoch	05	Train Loss	2.3399	Acc	10.25%		Test Loss	2.3501	Acc	9.80%
Epoch	06	Train Loss	2.3395	Acc	10.20%		Test Loss	2.3457	Acc	10.28%
Epoch	07	Train Loss	2.3384	Acc	10.22%		Test Loss	2.3221	Acc	9.82%
Epoch	08	Train Loss	2.3426	Acc	10.03%		Test Loss	2.3290	Acc	9.74%
Epoch	09	Train Loss	2.3406	Acc	10.33%		Test Loss	2.3643	Acc	9.74%
Epoch	10	Train Loss	2.3405	Acc	10.22%		Test Loss	2.3168	Acc	11.35%
Epoch	11	Train Loss	2.3389	Acc	10.25%		Test Loss	2.3176	Acc	11.35%
Epoch	12	Train Loss	2.3444	Acc	10.27%		Test Loss	2.3412	Acc	9.74%
Epoch	13	Train Loss	2.3416	Acc	10.03%		Test Loss	2.3422	Acc	9.80%
Epoch	14	Train Loss	2.3394	Acc	10.22%		Test Loss	2.3536	Acc	10.10%
Epoch	15	Train Loss	2.3388	Acc	10.30%		Test Loss	2.3536	Acc	10.32%
Epoch	16	Train Loss	2.3420	Acc	10.31%		Test Loss	2.3258	Acc	9.82%
Epoch	17	Train Loss	2.3410	Acc	10.38%		Test Loss	2.3331	Acc	11.35%
Epoch	18	Train Loss	2.3403	Acc	10.30%		Test Loss	2.3527	Acc	11.35%
Epoch	19	Train Loss	2.3392	Acc	10.18%		Test Loss	2.3767	Acc	9.82%
Epoch	20	Train Loss	2.3410	Acc	10.48%		Test Loss	2.3394	Acc	10.09%
Epoch	21	Train Loss	2.3408	Acc	9.98%		Test Loss	2.3222	Acc	11.35%
Epoch	22	Train Loss	2.3434	Acc	10.11%		Test Loss	2.3409	Acc	10.32%
Epoch	23	Train Loss	2.3434	Acc	10.22%		Test Loss	2.3340	Acc	9.58%
Epoch	24	Train Loss	2.3391	Acc	10.41%		Test Loss	2.3304	Acc	8.92%
Epoch	25	Train Loss	2.3398	Acc	9.81%		Test Loss	2.3388	Acc	9.82%
Epoch	26	Train Loss	2.3405	Acc	10.05%		Test Loss	2.3275	Acc	11.35%
Epoch	27	Train Loss	2.3401	Acc	10.04%		Test Loss	2.3304	Acc	10.09%
Epoch	28	Train Loss	2.3380	Acc	10.16%		Test Loss	2.3779	Acc	11.35%
Epoch	29	Train Loss	2.3393	Acc	10.18%		Test Loss	2.3491	Acc	8.92%
Epoch	30	Train Loss	2.3446	Acc	10.13%		Test Loss	2.3521	Acc	10.28%

- Training accuracy ~ 10%, Test accuracy ~ 20%
- Behavior: no learning. Loss ~ 2.34 constant
- Reason: the step size is too large; gradients overshoot the optimum, so weights never converge

## → 0.1 Learning Rate:

Epoch	Train Loss	Acc	Test Loss	Acc
Epoch 01	Train Loss 1.5334	Acc 40.35%	Test Loss 1.4108	Acc 48.18%
Epoch 02	Train Loss 1.4038	Acc 46.55%	Test Loss 1.5028	Acc 41.69%
Epoch 03	Train Loss 1.4454	Acc 41.22%	Test Loss 1.5707	Acc 37.36%
Epoch 04	Train Loss 1.4065	Acc 42.46%	Test Loss 1.3452	Acc 44.92%
Epoch 05	Train Loss 1.4601	Acc 40.21%	Test Loss 1.4619	Acc 43.11%
Epoch 06	Train Loss 1.5027	Acc 39.61%	Test Loss 1.6584	Acc 32.99%
Epoch 07	Train Loss 1.4806	Acc 39.52%	Test Loss 1.7789	Acc 28.99%
Epoch 08	Train Loss 1.5346	Acc 36.18%	Test Loss 1.5447	Acc 32.94%
Epoch 09	Train Loss 1.5553	Acc 33.87%	Test Loss 1.5642	Acc 29.99%
Epoch 10	Train Loss 1.5474	Acc 34.88%	Test Loss 1.5131	Acc 34.20%
Epoch 11	Train Loss 1.5124	Acc 36.39%	Test Loss 1.5476	Acc 37.20%
Epoch 12	Train Loss 1.5078	Acc 36.55%	Test Loss 1.4701	Acc 38.75%
Epoch 13	Train Loss 1.4716	Acc 37.89%	Test Loss 1.4949	Acc 41.88%
Epoch 14	Train Loss 1.4731	Acc 37.61%	Test Loss 2.2073	Acc 28.84%
Epoch 15	Train Loss 1.5482	Acc 35.97%	Test Loss 1.4586	Acc 40.28%
Epoch 16	Train Loss 1.4776	Acc 36.29%	Test Loss 1.4665	Acc 37.08%
Epoch 17	Train Loss 1.4815	Acc 36.35%	Test Loss 1.5287	Acc 34.06%
Epoch 18	Train Loss 1.4849	Acc 36.33%	Test Loss 1.4964	Acc 34.76%
Epoch 19	Train Loss 1.5196	Acc 35.06%	Test Loss 1.4924	Acc 36.22%
Epoch 20	Train Loss 1.5155	Acc 35.34%	Test Loss 1.4695	Acc 38.95%
Epoch 21	Train Loss 1.5022	Acc 35.45%	Test Loss 1.5045	Acc 36.87%
Epoch 22	Train Loss 1.5822	Acc 35.31%	Test Loss 1.5565	Acc 35.25%
Epoch 23	Train Loss 1.5044	Acc 35.63%	Test Loss 1.4681	Acc 38.44%
Epoch 24	Train Loss 1.4937	Acc 35.93%	Test Loss 1.4771	Acc 36.29%
Epoch 25	Train Loss 1.4973	Acc 35.71%	Test Loss 1.5212	Acc 31.47%
Epoch 26	Train Loss 1.4860	Acc 35.86%	Test Loss 1.4853	Acc 34.29%
Epoch 27	Train Loss 1.4977	Acc 35.55%	Test Loss 1.6560	Acc 33.58%
Epoch 28	Train Loss 1.5498	Acc 34.62%	Test Loss 1.4879	Acc 36.52%
Epoch 29	Train Loss 1.4935	Acc 35.44%	Test Loss 1.4994	Acc 34.98%
Epoch 30	Train Loss 1.5022	Acc 35.65%	Test Loss 1.5266	Acc 36.31%

- Training accuracy ~ 35 - 46%, Test Accuracy ~ 30 - 48%
- Behavior: model started to learn but accuracy fluctuated and plateaued early.
- Reason: still relatively high LR → unstable updates and oscillation around the optimum

## → 0.01 Learning Rate:

Epoch	Train Loss	Acc	Test Loss	Acc
Epoch 01	Train Loss 0.7331	Acc 77.42%	Test Loss 0.4884	Acc 86.81%
Epoch 02	Train Loss 0.4133	Acc 88.74%	Test Loss 0.4383	Acc 88.59%
Epoch 03	Train Loss 0.3657	Acc 90.09%	Test Loss 0.3715	Acc 90.02%
Epoch 04	Train Loss 0.3495	Acc 90.55%	Test Loss 0.3885	Acc 89.62%
Epoch 05	Train Loss 0.3360	Acc 90.87%	Test Loss 0.3601	Acc 90.01%
Epoch 06	Train Loss 0.3323	Acc 91.03%	Test Loss 0.3630	Acc 90.23%
Epoch 07	Train Loss 0.3300	Acc 91.13%	Test Loss 0.3588	Acc 90.37%
Epoch 08	Train Loss 0.3232	Acc 91.20%	Test Loss 0.3701	Acc 89.49%
Epoch 09	Train Loss 0.3173	Acc 91.44%	Test Loss 0.3736	Acc 90.05%
Epoch 10	Train Loss 0.3145	Acc 91.50%	Test Loss 0.3738	Acc 90.19%
Epoch 11	Train Loss 0.3162	Acc 91.39%	Test Loss 0.3479	Acc 90.73%
Epoch 12	Train Loss 0.3107	Acc 91.48%	Test Loss 0.3405	Acc 90.67%
Epoch 13	Train Loss 0.3074	Acc 91.59%	Test Loss 0.3502	Acc 90.83%
Epoch 14	Train Loss 0.3047	Acc 91.76%	Test Loss 0.3625	Acc 90.56%
Epoch 15	Train Loss 0.3054	Acc 91.73%	Test Loss 0.3380	Acc 90.78%
Epoch 16	Train Loss 0.3049	Acc 91.77%	Test Loss 0.3534	Acc 91.01%
Epoch 17	Train Loss 0.3041	Acc 91.72%	Test Loss 0.3530	Acc 90.81%
Epoch 18	Train Loss 0.3017	Acc 91.81%	Test Loss 0.3519	Acc 90.85%
Epoch 19	Train Loss 0.3034	Acc 91.78%	Test Loss 0.3720	Acc 90.24%
Epoch 20	Train Loss 0.3000	Acc 91.80%	Test Loss 0.3692	Acc 90.52%
Epoch 21	Train Loss 0.2972	Acc 91.90%	Test Loss 0.3693	Acc 90.11%
Epoch 22	Train Loss 0.3007	Acc 91.81%	Test Loss 0.3720	Acc 90.43%
Epoch 23	Train Loss 0.2954	Acc 92.03%	Test Loss 0.3523	Acc 90.93%
Epoch 24	Train Loss 0.2961	Acc 91.90%	Test Loss 0.3492	Acc 91.14%
Epoch 25	Train Loss 0.2960	Acc 92.08%	Test Loss 0.3586	Acc 90.75%
Epoch 26	Train Loss 0.2930	Acc 92.02%	Test Loss 0.3607	Acc 90.76%
Epoch 27	Train Loss 0.2940	Acc 92.08%	Test Loss 0.3542	Acc 91.16%
Epoch 28	Train Loss 0.2925	Acc 92.08%	Test Loss 0.3698	Acc 90.66%
Epoch 29	Train Loss 0.2924	Acc 92.17%	Test Loss 0.3611	Acc 90.81%
Epoch 30	Train Loss 0.2926	Acc 92.00%	Test Loss 0.3807	Acc 90.06%

- Training accuracy: went up from 77% to 92%, Test accuracy: increased to ~91%
- Behavior: stable convergence with smooth loss reduction from 0.73 → 0.29
- Reason: LR = 0.01 provided a balanced trade-off between speed and stability, allowing the model to learn digit patterns effectively.

→ 0.001 Learning Rate:

Epoch	Train Loss	Acc	Test Loss	Acc
Epoch 01	Train Loss 1.2107	Acc 59.89%	Test Loss 0.7440	Acc 77.49%
Epoch 02	Train Loss 0.6825	Acc 79.33%	Test Loss 0.6301	Acc 81.50%
Epoch 03	Train Loss 0.6021	Acc 82.20%	Test Loss 0.5738	Acc 83.50%
Epoch 04	Train Loss 0.5459	Acc 84.11%	Test Loss 0.5329	Acc 84.55%
Epoch 05	Train Loss 0.4970	Acc 85.53%	Test Loss 0.4896	Acc 86.05%
Epoch 06	Train Loss 0.4580	Acc 86.83%	Test Loss 0.4545	Acc 87.03%
Epoch 07	Train Loss 0.4288	Acc 87.70%	Test Loss 0.4276	Acc 87.80%
Epoch 08	Train Loss 0.4072	Acc 88.41%	Test Loss 0.4140	Acc 88.26%
Epoch 09	Train Loss 0.3910	Acc 88.97%	Test Loss 0.3970	Acc 88.99%
Epoch 10	Train Loss 0.3790	Acc 89.31%	Test Loss 0.3871	Acc 89.24%
Epoch 11	Train Loss 0.3678	Acc 89.68%	Test Loss 0.3822	Acc 89.34%
Epoch 12	Train Loss 0.3601	Acc 89.91%	Test Loss 0.3676	Acc 89.95%
Epoch 13	Train Loss 0.3526	Acc 90.16%	Test Loss 0.3649	Acc 90.04%
Epoch 14	Train Loss 0.3458	Acc 90.39%	Test Loss 0.3600	Acc 90.14%
Epoch 15	Train Loss 0.3409	Acc 90.51%	Test Loss 0.3576	Acc 90.33%
Epoch 16	Train Loss 0.3359	Acc 90.70%	Test Loss 0.3555	Acc 90.27%
Epoch 17	Train Loss 0.3318	Acc 90.79%	Test Loss 0.3555	Acc 90.39%
Epoch 18	Train Loss 0.3287	Acc 90.87%	Test Loss 0.3546	Acc 90.46%
Epoch 19	Train Loss 0.3261	Acc 90.97%	Test Loss 0.3462	Acc 90.74%
Epoch 20	Train Loss 0.3230	Acc 91.02%	Test Loss 0.3491	Acc 90.44%
Epoch 21	Train Loss 0.3207	Acc 91.17%	Test Loss 0.3513	Acc 90.66%
Epoch 22	Train Loss 0.3184	Acc 91.23%	Test Loss 0.3492	Acc 90.44%
Epoch 23	Train Loss 0.3159	Acc 91.25%	Test Loss 0.3443	Acc 90.54%
Epoch 24	Train Loss 0.3138	Acc 91.32%	Test Loss 0.3457	Acc 90.51%
Epoch 25	Train Loss 0.3127	Acc 91.38%	Test Loss 0.3476	Acc 90.67%
Epoch 26	Train Loss 0.3113	Acc 91.43%	Test Loss 0.3444	Acc 90.88%
Epoch 27	Train Loss 0.3092	Acc 91.53%	Test Loss 0.3456	Acc 90.90%
Epoch 28	Train Loss 0.3084	Acc 91.55%	Test Loss 0.3459	Acc 90.74%
Epoch 29	Train Loss 0.3061	Acc 91.55%	Test Loss 0.3458	Acc 90.74%
Epoch 30	Train Loss 0.3057	Acc 91.63%	Test Loss 0.3468	Acc 90.83%

- Training accuracy ~ 59% → 91%, Test accuracy ~ 77% → 90%
- Behavior: Learning is steady but slower than LR = 0.01
- Reason: updates are small; convergence is stable but requires more epochs to reach similar accuracy.

## Task 6

### Architecture:

- Input: 28x28 grayscale image → flatten to 784 features
- Hidden 1: Linear(784 → 8) + ReLU
- Hidden 2: Linear(8 → 4) + ReLU
- Output: Linear(4 → 10) → logits for digits 0..9
- Loss/activation: nn.CrossEntropyLoss (applies softmax internally to logits)

### Best performance so far:

- Test accuracy: ~91%
- Achieved with:
  - BATCH\_SIZE = 64
  - EPOCHS = 30
  - LR = 0.01
  - OPTIMIZER = "adam"
  - SEED = 42
  - LAYERS = [8, 4]

→ With this tiny two-hidden-layer MLP, LR = 0.01 struck the best balance of stability and speed, driving test accuracy to ~91%, whereas larger LR values were unstable and smaller LR (0.001) converged more slowly.

## Task 7

E456ComputerAssignment3/mlp.py			
Epoch 01	Train Loss 0.3729 Acc 88.96%	Test Loss 0.2515 Acc 92.63%	
Epoch 02	Train Loss 0.2384 Acc 93.09%	Test Loss 0.2283 Acc 93.49%	
Epoch 03	Train Loss 0.2071 Acc 94.01%	Test Loss 0.2296 Acc 93.81%	
Epoch 04	Train Loss 0.1972 Acc 94.24%	Test Loss 0.2305 Acc 93.38%	
Epoch 05	Train Loss 0.1857 Acc 94.57%	Test Loss 0.2257 Acc 93.68%	
Epoch 06	Train Loss 0.1812 Acc 94.69%	Test Loss 0.2050 Acc 94.22%	
Epoch 07	Train Loss 0.1745 Acc 94.89%	Test Loss 0.2257 Acc 93.80%	
Epoch 08	Train Loss 0.1681 Acc 95.11%	Test Loss 0.2175 Acc 93.85%	
Epoch 09	Train Loss 0.1679 Acc 95.06%	Test Loss 0.2104 Acc 94.33%	
Epoch 10	Train Loss 0.1624 Acc 95.25%	Test Loss 0.2123 Acc 94.28%	

Architecture:  $784 \rightarrow 16 \rightarrow 8 \rightarrow 10$  (ReLU activations)

Hyperparameters: EPOCHS = 10, LR = 0.01, BATCH\_SIZE = 64, OPTIMIZER = adam

Doubling the number of neurons in each hidden layer improved both training and test accuracy from ~91% (in Task 5) to ~ 94%, showing that the model benefited from increased capacity. The lower losses and smoother convergence indicate the network can now capture more complex digit patterns without overfitting, especially with only 10 epochs and a moderate learning rate.

## Task 8

E456ComputerAssignment3/mlp.py			
Epoch 01	Train Loss 0.3893 Acc 88.78%	Test Loss 0.2085 Acc 93.57%	
Epoch 02	Train Loss 0.1652 Acc 95.11%	Test Loss 0.1361 Acc 96.00%	
Epoch 03	Train Loss 0.1151 Acc 96.49%	Test Loss 0.1039 Acc 96.90%	
Epoch 04	Train Loss 0.0830 Acc 97.48%	Test Loss 0.0981 Acc 96.92%	
Epoch 05	Train Loss 0.0658 Acc 97.92%	Test Loss 0.0998 Acc 97.05%	
Epoch 06	Train Loss 0.0536 Acc 98.33%	Test Loss 0.0839 Acc 97.52%	
Epoch 07	Train Loss 0.0449 Acc 98.57%	Test Loss 0.1001 Acc 97.13%	
Epoch 08	Train Loss 0.0349 Acc 98.89%	Test Loss 0.0807 Acc 97.67%	
Epoch 09	Train Loss 0.0318 Acc 98.98%	Test Loss 0.0867 Acc 97.60%	
Epoch 10	Train Loss 0.0261 Acc 99.11%	Test Loss 0.1063 Acc 97.23%	
Epoch 11	Train Loss 0.0219 Acc 99.29%	Test Loss 0.1079 Acc 97.31%	
Epoch 12	Train Loss 0.0212 Acc 99.29%	Test Loss 0.0965 Acc 97.65%	
Epoch 13	Train Loss 0.0174 Acc 99.42%	Test Loss 0.1079 Acc 97.49%	
Epoch 14	Train Loss 0.0142 Acc 99.54%	Test Loss 0.1112 Acc 97.70%	
Epoch 15	Train Loss 0.0173 Acc 99.41%	Test Loss 0.1325 Acc 97.33%	

Architecture:  $784 \rightarrow 128 \rightarrow 64 \rightarrow 32 \rightarrow 10$  (ReLU activations)

Hyperparameters:

- BATCH\_SIZE = 64
- EPOCHS = 15
- LR = 0.001

Adding a third hidden layer and significantly increasing the number of neurons allowed the network to model much more complex relationships between pixels. The model now converges smoothly with very low loss and high test accuracy (~97.7%), outperforming all previous architectures.

The smaller learning rate (0.001) provided stable, fine-grained updates, preventing the instability seen at higher learning rates while ensuing fast convergence. The accuracy plateau around 97-98% suggests the model is approaching the limit of what a simple fully-connected MLP can achieve on MNIST without convolutional layers.

```
# ===== Hyperparameters =====
LAYERS    = [128, 64, 32]    # Hidden layer sizes.
BATCH_SIZE = 64            # Batch size
EPOCHS    = 15              # Number of training epochs
LR        = 0.001           # Learning rate
OPTIMIZER = "adam"         # Choose "adam" or "sgd"
STUDENT_ID = "907394064"
```

This seems to be the best architecture so far, achieving 97.76% test accuracy; which is the highest among all other experiments.

## Part B: CNN

E456ComputerAssignment3/cnn.py		
Epoch 01	Train Loss 0.2060 Acc 93.46%	Test Loss 0.0421 Acc 98.52%
Epoch 02	Train Loss 0.0523 Acc 98.39%	Test Loss 0.0299 Acc 99.06%
Epoch 03	Train Loss 0.0372 Acc 98.85%	Test Loss 0.0259 Acc 99.24%
Epoch 04	Train Loss 0.0292 Acc 99.05%	Test Loss 0.0263 Acc 99.10%
Epoch 05	Train Loss 0.0228 Acc 99.29%	Test Loss 0.0224 Acc 99.35%
Epoch 06	Train Loss 0.0187 Acc 99.39%	Test Loss 0.0266 Acc 99.15%
Epoch 07	Train Loss 0.0179 Acc 99.45%	Test Loss 0.0244 Acc 99.17%
Epoch 08	Train Loss 0.0138 Acc 99.54%	Test Loss 0.0227 Acc 99.34%
Epoch 09	Train Loss 0.0114 Acc 99.60%	Test Loss 0.0236 Acc 99.25%
Epoch 10	Train Loss 0.0108 Acc 99.63%	Test Loss 0.0241 Acc 99.24%
Epoch 11	Train Loss 0.0081 Acc 99.74%	Test Loss 0.0286 Acc 99.30%
Epoch 12	Train Loss 0.0099 Acc 99.64%	Test Loss 0.0245 Acc 99.34%

```
# Tunable parameters (kept the same as your settings)
BATCH_SIZE = 128
EPOCHS     = 12
LR          = 0.001
OPTIMIZER   = "adam"
DROPOUT_P   = 0.3
SEED        = 42
```

- Test accuracy crosses 98% and reaches a peak of ~99.35% at epoch 5.
- The final few epochs stay around 99.24-99.34% with low test losses (~0.022-0.029)
- The training loss decreases smoothly from 0.206 → 0.010 → 0.009, and training accuracy rises.

### Why this works:

- Adam + LR=1e-3 gives stable, fast convergence for this CNN on MNIST.
- Batch size 128 balances gradient quality with throughput.
- Dropout 0.3 in the classifier head adds mild regularization; no BatchNorm needed to surpass 98%.

With the above settings, the model consistently exceeds the requirement, achieving ~99.35% test accuracy while remaining stable (no overfitting signs). An early-stopping cutoff around epochs 5–8 would reach the same accuracy with less training time.